AADITYA K. SINGH

+44 7838 996940 (UK) | +1 703-731-2036 (US) | aaditya.singh.21@ucl.ac.uk | https://aadityasingh.github.io/

EDUCATION

University College London   **London, UK**
Gatsby Computational Neuroscience Unit, Ph.D. Student   *Sep. 2021 - Present*

Massachusetts Institute of Technology   **Cambridge, MA**
GPA: 5.0/5.0   *Sep. 2017 - Jun. 2021*
- M.Eng. and B.Sc. in Computer Science and Engineering, B.Sc. in Brain and Cognitive Sciences

PUBLICATIONS

Llama team, AI@Meta (Contributors: A. K. Singh, ...). **The Llama 3 herd of models**. 2024. ⓞ

A. K. Singh, T. Moskovitz, F. Hill, S. C. Y. Chan[†], A. M. Saxe[†]. **What needs to go right for an induction head? A mechanistic study of in-context learning circuits and their formation.** *ICML 2024* (Spotlight). https://arxiv.org/abs/2404.07129. ⓞ

A. K. Singh, DJ Strouse. **Tokenization counts: the impact of tokenization on arithmetic in frontier LLMs.** *In submission.* https://arxiv.org/abs/2402.14903. ⓞ

A. K. Singh, Y. Yang, K. Tirumala, M. Elhoushi, A. S. Morcos. **Brevity is the soul of wit: Pruning long files for code generation.** *ICML DMLR workshop (2024).* https://arxiv.org/abs/2407.00434

A. K. Singh*, S. C. Y. Chan*, T. Moskovitz, E. Grant, A. M. Saxe[†], F. Hill[†]. **The transient nature of emergent in-context learning in transformers.** *NeurIPS 2023.* https://arxiv.org/abs/2311.08360. ⓞ

T. Moskovitz, A. K. Singh, DJ Strouse, T. Sandholm, R. Salakhutdinov, A. D. Dragan, S. McAleer. **Confronting reward model overoptimization with constrained RLHF.** *ICLR 2024* (Spotlight). https://arxiv.org/abs/2310.04373

Y. Yang, A. K. Singh, M. Elhoushi, A. Mahmoud, K. Tirumala, F. Gloeckle, B. Roziere. C. Wu, A. S. Morcos, N. Ardalani. **Decoding data quality via synthetic corruptions: embedding-guided pruning of code data**. *NeurIPS ENSLP workshop 2023* (Oral spotlight). https://arxiv.org/abs/2312.02418

A. K. Singh, D. Ding, A.M. Saxe, F. Hill, A. K. Lampinen. **Know your audience: specializing grounded language models with the game of Dixit.** *EACL 2023.* https://arxiv.org/abs/2206.08349

S. C. Y. Chan, A. Santoro, A. K. Lampinen, J. X. Wang, A. K. Singh, P. H. Richemond, J. McClelland, F. Hill. **Data distributional properties drive emergent in-context learning in transformers.** *NeurIPS 2022* (Oral). https://arxiv.org/abs/2205.05055. ⓞ

RESEARCH EXPERIENCE

Gatsby Computational Neuroscience Unit   **London, UK**
PhD Student   *Sep. 2021 - Present*
- Co-supervised by Prof. Andrew Saxe and Dr. Felix Hill.
- Researching the transience of emergent few-shot learning in transformers from an empirical, mechanistic, and theoretical lens.
- Exploring the effects of number tokenization on math reasoning.
- Investigating how humans and RL agents learn abstract concepts in Sudoku-esque Nikoli puzzles. Humans transfer high-level concepts across such puzzles (with different state spaces), and our goal is to discover RL algorithms that do the same.
- Lead student-faculty representative: Aggregate and voice student concerns to faculty, and work to solve them.

Meta AI Research   **Menlo Park, CA**
Research Scientist Intern   *Jun. 2023 - Dec. 2023*
- Part of the Data Curation team in FAIR Labs, led by Dr. Ari S. Morcos, and then the Llama 3 team.
- Developed embedding-based and heuristic methods for pruning code data. NeurIPS and ICML workshop papers.
- Contributed to LLaMa 3 efforts, at all parts of the pipeline: Data preprocessing (for math reasoning), Scaling Laws, Evaluations

DeepMind   **London, UK**
Research Engineering Intern   *Jun. 2021 - Apr. 2022*
- Part of the Grounded Language team led by Dr. Felix Hill and Dr. Jane X. Wang.
- Led a 5-person project on finetuning grounded language models without direct supervision. EACL 2023 paper.
- Contributed to a larger project characterizing properties of data crucial for emergent few-shot learning. NeurIPS 2022 paper.
- Presented final results directly at organization-wide Research Lead meeting (RPM).

MIT InfoLab   **Cambridge, MA**
Research Assistant   *Jun. 2020 - Jun. 2021*
- Part of MIT CSAIL and MIT CBMM. Supervised by Prof. Boris Katz and collaborated with Prof. Ila Fiete.
- Led a project on bio-inspired deep attentional modulation for few-shot learning in object recognition. Master's Thesis (Grade: A)
- Core contributor to a project relating human intracranial recordings to language features (e.g., part of speech). In submission.
- Investigated Expander Hopfield Networks and their applicability for few-shot learning. Presented at lab meeting.
- Co-mentored three undergrad students and one high school student.

Citadel Securities **Chicago, IL**
Quantitative Research Intern *Jun. 2019 - Aug. 2019*
- Developed novel model selection techniques that led to 10% improvement and a 5x speed-up in equities alpha generation.
- Pioneered genetic algorithm methods for continuous blackbox optimization that led to 4% improvement over baseline alphas.
- Implemented and optimized Hidden Markov Model variants for prediction.
- Summarized work in 5 internal reports. Helped port algorithms to production.

Orbital Insight **Boston, MA**
Computer Vision Research Intern *Jan. 2019*
- Implemented a Discriminative Autoencoder for unsupervised change detection of satellite images using Tensorflow.
- Developed a weakly supervised model and experimented with different architectures, loss metrics and hyperparameters.
- Presented initial results to the Computer Vision team and VP of Tech Research.

Undergraduate Research **Cambridge, MA**
- Compositional ensemble learning. HackMIT Best Use of Data prize (2019). ☊
- Cross-movement art generation with a variational autoencoder. MIT College of Computing Launch Poster Session. ☊
- Physical-based audiovisual simulation in automatic online perception. Class project and runnable MTurk experiment. ☊
- Comparisons to Bayesian neural networks for weight pruning. Class project. ☊

Naval Research Labs **Washington, DC**
SEAP Researcher, Lab for Computational Physics and Fluid Dynamics *Jun. 2017 - Aug. 2017*

Metron, Inc. **Reston, VA**
Analyst Intern, Category Theory *Jun. 2016 - Aug. 2016*

ENGINEERING EXPERIENCE

FeatureX **Boston, MA**
Software Engineering Intern *Jun. 2018 - Aug. 2018*

Bublup **McLean, VA**
Software Engineering Intern *Jun. 2015 - Aug. 2015*

TEACHING EXPERIENCE

UCL Gatsby PhD courses **Sep. 2022 - Present**
- New student "last black box" bootcamp, Systems and Theoretical Neuroscience, Probabilistic and Unsupervised Learning
- Created new problem sets on: expectation maximization, variational inference, expectation propagation, deep linear networks

MIT 6.867 Graduate Machine Learning **Sep. 2019 - Dec. 2019**
- Ideated and wrote exam questions, comprising a third of the midterm and half of the final.
- Led recitation sessions for over 40 students and mentored eight project teams. Rated 7/7 by students in the course evaluation.

INVITED TALKS

- Gatsby Foundation Tricentre Meeting (Jun. 2024) - Learning dynamics of in-context learning in transformers
- Sainsbury Wellcome Center Symposium (Mar. 2024) - Neuroscience on neural networks
- Gatsby Foundation Scientific Advisory Board (Feb. 2024) - Learning dynamics of in-context learning in transformers
- MIT, Fiete Lab (Jan. 2024) - Learning dynamics of in-context learning circuits
- University of Witwatersrand, NLP Advanced Topics lecture (Oct. 2023) - the industrial LLM pipeline + mechanistic interpretability
- DeepMind Analysis Group (May 2023) - Emergence and transience of in-context learning in transformers
- Summerfield Lab (Jan. 2023) - Concept formation in puzzles

AWARDS

- Third place poster, Citadel PhD Summit (2024)
- Hertz Fellowship Finalist (2021)
- MIT Nominee for Marshall and Rhodes Scholarships (2021)
- MIT Brain and Cognitive Sciences Academic Achievement Award (2020)
- MIT CS+HASS Undergraduate Research and Innovation Scholar (2019-2020)
- MIT SHASS Burchard Scholar (2020)
- Top 16 in the US, MIT Battlecode (2018)
- 2nd place team in the US, National Science Bowl (2017)
- US National Olympiads: Silver Medalist (Physics, 2017), 14th in the Nation (Computational Linguistics, 2017), Top 50 in the Nation (Chemistry, 2017), USA(J)MO Qualifier (Math, 2014-2017), Platinum Divison (Computing, 2017)

PROGRAMMING SKILLS

- Proficient: Python, Numpy, JAX, PyTorch, Tensorflow, Java, MATLAB
- Competent: Bash, pyspark, LaTeX, ffmpeg, JavaScript, C, Ruby, Rails, HTML5, JQuery, CSS

RELEVANT COURSEWORK

Artificial Intelligence

Probabilistic & Unsupervised Learning, Kernels, Reinforcement Learning, Bayesian Inference, Statistical Learning Theory, Machine Learning, Natural Language Processing, Artificial Intelligence*

Math

Theory of Computation, Stochastic Processes, Matrix Methods for Machine Learning, Differential Equations*, Complex Analysis*, Probability Theory*, Linear Algebra*

Neuroscience

Systems & Theoretical Neuroscience, Neural Circuits for Cognition, Computational Cognitive Science, Systems Neuroscience Lab, Molecular & Cellular Neurobiology, Organic Chemistry

Software

Computer Security, Computer Systems Engineering, Elements of Software Construction, Design & Analysis of Algorithms, Parallel Computing*

indicates UG-level classes taken at Thomas Jefferson High School for Sci/Tech, 2015-2017